



How the policy learns: hPPO seed 10 from (1.0, 0.0) at evaluation checkpoints


10k steps: did not reach the goal (stopped at $x = 3.8$ m)




133k steps: did not reach the goal (stopped at $x = 2.8$ m)




256k steps: did not reach the goal (stopped at x =




389k steps: did not reach the goal (stopped at $x = 6.0$ m)



522k steps: did not reach the goal (stopped at $x = 6$)



1024k steps (end): did not reach the goal (stopped at $x = 6.6$ m)



— hPPO, seed 10

```
-- oracle
```

✗ wall contact

○ start