



How the policy learns: hPPO seed 11 from (1.0, 0.0) at evaluation checkpoints


10k steps: 99 steps




20k steps: 82 steps



31k steps (first solve): 80 steps



205k steps (end): 78 steps



— hPPO, seed 11 - - - oracle × wall contact ○ start